

# eda

December 10, 2025

## 1 Ursprung der Daten

Titanic Datensatz von [kaggle](#) heruntergeladen am 21.11.2025 um 19:22 Uhr.

## 2 Imports

```
[32]: %matplotlib inline

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import joblib
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve,roc_auc_score
```

## 3 Daten laden und erste Eindrücke sammeln

```
[33]: df = pd.read_csv("./data/Titanic-Dataset.csv")
```

### 3.1 Head

```
[34]: df.head()
```

```
[34]:   PassengerId  Survived  Pclass \
0            1         0      3
1            2         1      1
2            3         1      3
3            4         1      1
4            5         0      3

                                         Name     Sex   Age  SibSp \
0           Braund, Mr. Owen Harris    male  22.0      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0      1
2                Heikkinen, Miss. Laina  female  26.0      0
3        Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0      1
```

```

4                               Allen, Mr. William Henry    male  35.0      0
Parch      Ticket     Fare Cabin Embarked
0         A/5 21171  7.2500   NaN     S
1            PC 17599  71.2833  C85     C
2        STON/O2. 3101282  7.9250   NaN     S
3            113803  53.1000  C123     S
4            373450  8.0500   NaN     S

```

### 3.2 Info

[35]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
---  --  
 0   PassengerId  891 non-null   int64  
 1   Survived     891 non-null   int64  
 2   Pclass       891 non-null   int64  
 3   Name         891 non-null   object 
 4   Sex          891 non-null   object 
 5   Age          714 non-null   float64 
 6   SibSp        891 non-null   int64  
 7   Parch        891 non-null   int64  
 8   Ticket       891 non-null   object 
 9   Fare          891 non-null   float64 
 10  Cabin         204 non-null   object 
 11  Embarked     889 non-null   object 
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

```

### 3.3 Describe

[36]: df.describe()

```

[36]:    PassengerId  Survived  Pclass      Age      SibSp \
count    891.000000  891.000000  891.000000  714.000000  891.000000
mean     446.000000   0.383838   2.308642  29.699118   0.523008
std      257.353842   0.486592   0.836071  14.526497   1.102743
min      1.000000   0.000000   1.000000   0.420000   0.000000
25%     223.500000   0.000000   2.000000  20.125000   0.000000
50%     446.000000   0.000000   3.000000  28.000000   0.000000
75%     668.500000   1.000000   3.000000  38.000000   1.000000
max     891.000000   1.000000   3.000000  80.000000   8.000000

```

Parch Fare

```
count    891.000000  891.000000
mean      0.381594   32.204208
std       0.806057   49.693429
min       0.000000   0.000000
25%      0.000000   7.910400
50%      0.000000  14.454200
75%      0.000000  31.000000
max       6.000000  512.329200
```

### 3.4 Unique Values

```
[37]: df.nunique().sort_values()
```

```
[37]: Survived      2
Sex           2
Pclass        3
Embarked      3
SibSp         7
Parch         7
Age           88
Cabin         147
Fare          248
Ticket        681
PassengerId   891
Name          891
dtype: int64
```

### 3.5 Erkenntnisse aus den ersten Eindrücken

- bei Cabin fehlen zu viele Daten.
- bei Ticket, PassengerId und Name gibt es zu viele unterschiedliche Daten
- bei Age gibt es 177 Null-Werte
- Sex und Embarked müssen encoded werden

## 4 Datenbereinigung

### 4.1 Cabin, Ticket, PassengerId und Name entfernen

```
[38]: df.drop(columns=["PassengerId", "Name", "Cabin", "Ticket"], axis=1, ↴
               inplace=True)
```

## 4.2 Null-Werte bei Age mit dem Durchschnitt auffüllen

### 4.2.1 Verteilung des Alters


```
[39]: sns.set(style="whitegrid")

plt.figure(figsize=(12, 6))

sns.histplot(df["Age"], bins=30, kde=True, color="steelblue")

plt.title("Verteilung der Altersstruktur (Titanic)", fontsize=16)
plt.xlabel("Alter", fontsize=14)
plt.ylabel("Anzahl", fontsize=14)

plt.show()
print(f"Durchschnittsalter {df['Age'].mean():.2f}")
print(f"Median-Alter {df['Age'].median():.2f}")
```



Durchschnittsalter 29.69911764705882

Median-Alter 28.0

Durchschnittsalter und Median sind ähnlich hoch, der Median wird nun aber bevorzugt verwendet, da dieser die Ausreißer kompensieren soll.

```
[40]: df['Age'] = df['Age'].fillna(df['Age'].median())
```

## 4.3 Encoding vom Sex

```
[41]: df['Sex'] = df['Sex'].map({'male': 0, 'female': 1})
```

## 4.4 Auffüllen und Encoding von Embarked (Abfahrtsort)

### 4.4.1 Null-Werte werden mit den am häufigsten vorkommendem Wert aufgefüllt

```
[42]: df['Embarked'] = df['Embarked'].fillna(df['Embarked'].mode()[0])
```

### 4.4.2 One-Hot-Encoding

Aus der Spalte ‘Embarked’ werden die binären Spalten ‘Embarked\_S’ und ‘Embarked\_Q’. Der Abfahrtsort C wird abgeleitet und trifft dann zu, wenn ‘Embarked\_S’ und ‘Embarked\_Q’ beide 0 sind.

```
[43]: df = pd.get_dummies(df, columns=['Embarked'], drop_first=True)
```

## 5 Daten nach der Bereinigung

### 5.1 Head

```
[44]: df.head()
```

```
[44]:   Survived  Pclass  Sex    Age  SibSp  Parch     Fare  Embarked_Q  Embarked_S
 0         0      3    0  22.0      1      0    7.2500    False      True
 1         1      1    1  38.0      1      0   71.2833    False     False
 2         1      3    1  26.0      0      0    7.9250    False      True
 3         1      1    1  35.0      1      0   53.1000    False      True
 4         0      3    0  35.0      0      0    8.0500    False      True
```

### 5.2 Info

```
[45]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype  
 ____ _-----_ _-----_ _-----_
 0   Survived    891 non-null    int64 
 1   Pclass       891 non-null    int64 
 2   Sex          891 non-null    int64 
 3   Age          891 non-null    float64
 4   SibSp        891 non-null    int64 
 5   Parch        891 non-null    int64 
 6   Fare          891 non-null    float64
 7   Embarked_Q   891 non-null    bool   
 8   Embarked_S   891 non-null    bool
```

```
dtypes: bool(2), float64(2), int64(5)
memory usage: 50.6 KB
```

## 6 Train-Test-Split

```
[46]: X = df.drop(columns=["Survived"])
y = df['Survived']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

## 7 Modell-Auswahl und Training

Bei dem Datensatz soll eine Klassifikation vorgenommen werden, daher eignet sich beispielsweise Logistische Regression

```
[47]: model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
```


```
[47]: LogisticRegression(max_iter=1000)
```

## 8 Metriken

```
[48]: y_pred = model.predict(X_test)
acc = accuracy_score(y_test, y_pred)
```


### 8.1 Confusion Matrix

```
[49]: cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title(f'Confusion Matrix\nAccuracy: {acc:.3f}')
plt.savefig('./data/confusion_matrix.png', bbox_inches='tight')
plt.show()
```



## 9 ROC


```
[50]: roc_auc = roc_auc_score(y_test, y_pred)
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc:.3f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title(f'ROC - Logistic Regression\nAccuracy: {acc:.3f}')
plt.legend(loc="lower right")
plt.savefig('./data/roc_curve.png', bbox_inches='tight')
plt.show()
```



## 9.1 Korrelationen

```
[51]: corr = df.corr()["Survived"].drop("Survived").sort_values()

plt.figure(figsize=(8, 5))
sns.barplot(x=corr.values, y=corr.index)
plt.title("Korrelationen mit dem Target")
plt.xlabel("Korrelationsstärke")
plt.ylabel("Feature")
plt.savefig('./data/correlations.png', bbox_inches='tight')
plt.show()
```



## 10 Modell speichern

```
[52]: joblib.dump(model, "./data/model.pkl")
joblib.dump(X.columns, "./data/feature_names.pkl")
```

```
[52]: ['./data/feature_names.pkl']
```